

Nicholas Pezarro

San Francisco, CA · n.pezarro@gmail.com · linkedin.com/in/nicholaspezarro · github.com/npezarro

Summary

Technical product leader who builds developer tools, not just specs. Daily Claude Code power user who identified product gaps and built compensating infrastructure: a behavioral governance framework propagated across 60+ repositories, an A/B evaluation framework measuring how instruction environments affect output quality, and autonomous agent orchestration with self-improving feedback loops. 10+ years of product experience spanning generative AI, developer tooling, ML personalization, and growth. One of two primary prompt engineers for LinkedIn Collaborative Articles, where I defined behavioral defaults through a multi-stage prompt pipeline, built a three-layer evaluation framework that became the product's quality standard, and shipped responsible AI quality gates, driving zero to 900K weekly SEO sessions in 7 months. Built LinkedIn Games from concept to 7.5M+ WAU. Named inventor on two US patents (20240273306, 20250337653). Have helped multiple non-technical users bootstrap their own Claude Code environments from scratch, demonstrating that the product's power can extend well beyond developers when the on-ramp is designed well.

Independent Technical Work

Claude Code Agent Ecosystem · github.com/npezarro - 2024 - Present

- Behavioral Governance Framework (agentGuidance): Identified that Claude Code lacked cross-project consistency and session-to-session learning; built a 5-layer instruction hierarchy (global behavioral defaults, environment constraints, per-repo CLAUDE.md, cross-session memory, synthesized wiki) with SessionStart hooks that propagate rule changes to every future session without manual intervention. 33 guidance files covering testing, deployment, security, debugging, and session lifecycle. Framework is reusable: multiple non-technical users have bootstrapped their own Claude Code environments using it as a starting point.
- Self-Improving Instruction System (Learning Agent): Observed that behavioral corrections made during sessions were consistently lost; built an hourly agent (368+ runs) that reviews cross-repo activity, detects uncaptured corrections, audits memory for multi-destination compliance, and stages documentation PRs. Treats instruction files as first-class artifacts requiring maintenance, creating a compounding feedback loop where every session makes the system better.
- A/B Evaluation Framework (claude-bakeoff): Built after observing that small prompt changes produced unpredictable behavioral cascades across agent sessions. Tests identical tasks under different instruction environments using LLM-as-judge scoring with structured rubrics (correctness, completeness, code quality, adherence). 80% correlation with human preference. Enables evidence-based decisions about instruction changes before deploying to production agents.
- Autonomous Development Agent (autonomousDev): Built a 30-minute cron agent (365+ production runs) that surveys 26+ repos for actionable work (build failures, security vulnerabilities, lint errors, dependency gaps), executes via Claude Code CLI governed by agentGuidance, and creates PRs for human review. Never auto-merges: deliberate safety decision after analyzing failure modes. 5 protected repos require manual approval.
- Cross-Instance Knowledge Sharing (ClaudeNet): Identified that context died between Claude Code sessions and across instances; built an async messaging system where separate CLIs exchange implementation knowledge via threaded conversations. Express.js + SQLite, 168 tests, deployed on production infrastructure.

- Behavioral Taxonomy (15 Specialist Profiles): Designed 15 specialist behavioral profiles (architect, critic, implementer, security, QA, PM, and others), each with personality parameters, expertise boundaries, deference rules, and append-only experience logs that accumulate knowledge across sessions. Profiles can be versioned and A/B tested via claude-bakeoff.
- Supporting Infrastructure: 10+ deployed web applications (Next.js/Prisma/PostgreSQL, Vite/React, Chrome extensions). Open-source browser agent (MV3 extension + CDP integration, 30+ commands) enabling AI CLI tools to control a live browser with trusted keyboard events. Discord-based orchestration platform with 20+ auto-managed channels, 14 task templates, multi-persona debate system, session pool management, and crash-resilient job dispatch. Two-tier error recovery system (fix-checker) that restarts first, then dispatches Claude to analyze failures and stage code fix PRs. Bidirectional health monitoring across WSL and GCP VM infrastructure.

Experience

LinkedIn, San Francisco, CA - AUGUST 2021 - PRESENT

Senior Product Manager (Product Lead, Games & GenAI)

- Generative AI Product Lead (Collaborative Articles): Defined behavioral defaults for LinkedIn's first large-scale LLM content product through a 4-turn prompt pipeline (Topic, Structure, Body, Title), with parameter tuning calibrated per generation stage and semantic constraints preventing hallucination and format violations. Maintained a behavioral taxonomy of model failure modes with programmatic detection and recovery prompts. Wrote LinkedIn's internal Prompt Engineering Best Practices playbook, scaling methodology across teams adopting generative AI.
- Three-Layer Evaluation Framework: Built a multi-layer eval system: automated quality scoring to detect format and content violations, human calibration to catch gaps invisible to automation (discovered 30% divergence between automated and human scores when articles lacked expert contributions), and real-world performance feedback validating choices against user outcomes. Architecture formalized in US Patent 20240273306.
- Responsible AI Quality Gate: Implemented a quality gate requiring 4+ expert contributions before AI content was surfaced to millions. Improved indexation 25% (66% to 82%) and drove +6.4% session lift, demonstrating that responsible AI constraints can strengthen product performance. Drove 900K weekly SEO sessions within 7 months from zero.
- Safety-vs-Utility Tradeoff Analysis (Anti-Scraping): Led quantitative analysis protecting 900M+ public profiles from AI-driven data scraping (2B+ monthly hits). Designed structured experiment framework testing 4 variants across 23M profiles. Achieved -95% scraping at only -0.6% WAU impact. Presented tradeoff framework directly to CEO and CPO. Protected \$200M-\$320M in enterprise bookings.
- Consumer Growth, LinkedIn Games (0-to-1): Part of a founding team of 3 PMs who scaled LinkedIn Games from concept to 7.5M+ WAU. Identified a 66-77% correlation between win-rate and D1R as the primary retention lever, then shipped targeted interventions that moved new player D1R from 17.44% to 19.75% and D7R from 10.64% to 12.37%.
- ML Content Recommendation & Ranking: Built EBR (Embedding-Based Retrieval) crosslinking system delivering +12% Content SEO lift. Led 5x organic session growth (2.73M to 13M sessions). Owned offline evaluation methodology, launch criteria, and post-launch monitoring.

Tophatter, San Francisco, CA · JUNE 2020 - JULY 2021

Product Lead

- Launched "Name Your Price" fixed-price feature from concept to 15% of total revenue in 75 days with a 3-engineer pod; +30% AOV with zero auction cannibalization. Led UX revamp delivering 2x CSAT improvement; elevated Google Play rating from 3.6 to 4.5.

Curology, San Francisco, CA · SEPTEMBER 2019 - MARCH 2020

Senior Product Manager

- Drove acquisition of 5,000+ users/day. Led identity verification implementation handling 70% of signups without friction. Restructured experimentation using statistical power analysis.

Humble Bundle, San Francisco, CA · DECEMBER 2016 - AUGUST 2019

Product Manager

- Launched new homepage driving 50% increase in average visit value. Owned fraud prevention leveraging Sift Science and Smyte ML scoring APIs, designing rule architectures with precision/recall tradeoffs, reducing dispute volume 69% and gross losses 71% YoY. Led emergency migration from Smyte to Sift Workflows after Twitter's acquisition shut down the service overnight.

Pocket Gems, San Francisco, CA · JULY 2015 - NOVEMBER 2016

Associate Product Manager

- Managed features and live events for War Dragons, a top-25-grossing mobile title; introduced a new event type that drove 100% increase in ARPDau and sustained 20-30% revenue uplifts.

Intellectual Property

US Patent Application 20240273306: "Generative Collaborative Publishing System", AI-generated professional content with human expert contributions, prompt pipeline architecture, and collaborative publishing mechanics.

US Patent Application 20250337653: "Artificial Intelligence Techniques for Connections Networking", ML-based entity grouping and member recommendation system.

Skills & Technical Expertise

Developer Tools & Ecosystem: CLI Tooling, Developer Experience (DX) Design, Extension/Plugin Architecture, Configuration Ecosystem Design, Ecosystem Evangelism (60+ public repos, technical blog)

AI Systems & Behavioral Design: Behavioral Default Design, Steerability Constraints, Prompt Pipeline Architecture, Parameter Tuning, Semantic Guardrails, Hallucination Prevention, Agent Behavioral Profiles, Rule Propagation, Failure Mode Taxonomy

Evaluation & Alignment: Multi-layer Eval Framework Design (Automated + Human + Real-World), LLM-as-Judge Scoring, A/B Instruction Environment Testing, Safety-vs-Utility Tradeoff Quantification, Human-in-the-Loop Quality Gates

Generative AI & LLM Products: Prompt Engineering (Multi-turn Chains, Few-shot, Chain-of-Thought), RAG & Embeddings, AI Coding Assistants (Claude Code CLI daily user), JavaScript/TypeScript, Python

ML & Growth: Content Recommendation Models, EBR, Ranking Signals, Offline/Online A/B Testing, Statistical Power Analysis, Cohort & Funnel Analysis

Education & Continuous Learning

Fast.ai - *Practical Deep Learning for Coders (2020)*

Reforge - *Advanced Growth Series, Retention + Engagement Deep Dive*

University of British Columbia, Vancouver, BC - *Bachelor of Commerce*